

- [Contact Us](#)
- [Log in](#)
 - [Monitor](#)
 - [Secure](#)
 - [Support](#)
- -
 - [日本](#)
 - [United States](#)
-

Search for:

[Products](#)[Use Cases](#)[Open Source](#)[Resources](#)[Company](#)[Start Free Trial](#)

- [Products](#)
 - [Sysdig Platform](#)
 - [Sysdig Secure](#)
 - [Sysdig Monitor](#)
 - [Getting Started](#)
 - [Pricing](#)
 - [Start Free Trial](#)
- [Use Cases](#)
 - [Kubernetes Security](#)
 - [Runtime Security](#)
 - [Image Scanning](#)
 - [IR & Forensics](#)
 - [Compliance](#)
 - [Kubernetes Monitoring](#)
 - [Full Stack Monitoring](#)
 - [Troubleshooting](#)
 - [Prometheus Monitoring](#)
 - [Dashboards](#)
- [Open Source](#)
 - [Sysdig Contributions](#)
 - [Enterprise Falco](#)
 - [Prometheus](#)
 - [sysdig Inspect](#)
- [Environments](#)
 - [AWS](#)
 - [Google Cloud](#)
 - [IBM Cloud](#)
 - [Microsoft Azure](#)
 - [OpenShift](#)
 - [Federal Agencies](#)
- [Resources](#)
 - [Library](#)
 - [Webinars](#)
 - [Case Studies](#)
 - [Blogs](#)
 - [Newsletters](#)
- [Company](#)
 - [About Us](#)
 - [Leadership](#)
 - [Trust Center](#)
 - [Careers](#)
 - [Events](#)
 - [Webinars](#)
 - [Request a Demo](#)

Updated Recommendations for You



- [Newsroom](#)
 - [Newsroom](#)
 - [Press Releases](#)
 - [News Coverage](#)
 - [Media Resources](#)

Search for:

[Start 30-day Free Trial](#)

[Contact Sales](#)

Follow us.



Copyright 2020 Sysdig, Inc. All Rights Reserved.

[Privacy Policy](#)



Blog Post

How to write a custom Kubernetes scheduler using your monitoring metrics

By Mateo Burillo
on November 3, 2017

LIVE WEBINAR: Getting Started with Prometheus Exporters - August 18, 2020 10am Pacific / 1pm Eastern

[Register Now](#)

[Watch Webinars Now](#)

This article covers the use case of creating a custom Kubernetes scheduler and implements an example using monitoring metrics to make scheduling decisions.

Updated Recommendations for You

^

UPDATE: There is a new and more complete implementation of the custom [Kubernetes scheduler using Golang](#).

The default Kubernetes scheduler does a fantastic job for most typical workloads. Starting from Kubernetes 1.6 [advanced scheduling features](#) like node or pod affinity, taints and tolerations allows you to configure several pod scheduling policies: in a specific set of nodes (node affinity/anti-affinity), close or far away from other running pods (pod affinity/anti-affinity), or just based on some tags that pods like or dislike (taints and tolerations).

But maybe you have some more specific requirements or would like to use higher level and dynamic application information to map your new pods to the physical nodes. Always striving for extensibility and flexibility, Kubernetes 1.6 introduced [multiple scheduler/custom scheduler](#) support as a beta feature.

What if you could use any of the metrics already present in your Kubernetes monitoring system to configure the behaviour of your pod scheduler?

How to write a custom #Kubernetes scheduler using your monitoring metrics

[Click to tweet](#)

The following is an example of a custom scheduler using metrics from our Kubernetes monitoring tool: [Sysdig Monitor](#). In Sysdig, all metrics are automatically tagged with Kubernetes metadata, so you can easily do advanced monitoring, alerting, troubleshooting and now, advanced scheduling too.

Coding this scheduler may be a lot simpler that you may imagine. Let's start with a simple example to give you some context and ideas. Say for example that you want to optimize the responsiveness that your users perceive, so you decide that you want to place new web server pods in the physical host that is scoring the best **HTTP response times** at that specific point in time.

Normally, as a prerequisite you would have to instrument your application, but Sysdig collects requests, errors and response times metrics for any application or service without any kind of code instrumentation. But if you wanted to write the scheduler based on the behavior of an internal application metric, Sysdig will get any statsd, JMX or [Prometheus metrics](#) for you automatically, awesome! Isn't it?

Configure your pods to use a custom Kubernetes scheduler

First, you need to configure your pods to use a custom scheduler:

```
1  apiVersion: v1
2  kind: ReplicationController
3  metadata:
4    name: nginx
5  spec:
6    replicas: 3
7    selector:
8      app: nginx
9    template:
10     metadata:
11       name: nginx
12       labels:
13         app: nginx
14     spec:
15       schedulerName: sysdigsched
16     containers:
17     - name: nginx
18       image: nginx
19       ports:
20       - name: http
21         containerPort: 80
```

nginx_scheduler.yaml hosted with ❤ by GitHub [view raw](#)

This is a very simple vanilla Nginx replicationController. Note that we added `schedulerName: sysdigsched` to the pod definition. Remember that this is a Kubernetes 1.6+ feature, so this config parameter will throw an error when using older versions.

If you push this replicationController to the cluster:

```
$ kubectl create -f nginxrc.yaml
replicationcontroller "nginx" created
```

Updated Recommendations for You

^

NAME	READY	STATUS	RESTARTS	AGE
nginx-84cnn	0/1	Pending	0	11s
nginx-ffldk	0/1	Pending	0	11s
nginx-jq5jk	0/1	Pending	0	11s

The pods will never leave the `Pending` state. They require a pod scheduler that doesn't yet exist.

Write your own Kubernetes scheduler!

We are going to use Python for this example. First, you will need to install both Kubernetes and [Sysdig python libraries](#):

```
pip install kubernetes sdcclient
```

Before jumping to the complete example, let's look at the (relevant sections of the) example code.

First, you need to import the Kubernetes and [Sysdig API](#) libraries and objects:

```
from kubernetes import client, config, watch
from sdcclient import SdcClient
```

Then, you can initialize the API objects. You will find the *Sysdig Monitor API Token* for your account in the *User Profile* section of your Sysdig Monitor configuration menu:

```
config.load_kube_config()
v1=client.CoreV1Api()
sdclient = SdcClient(<your_sysdig_token>)
sysdig_metric = "net.http.request.time"
metrics = [{ "id": sysdig_metric, "aggregations": { "time": "timeAvg", "group": "avg" } }]
</your_sysdig_token>
```

For this example, we are only going to use the `net.http.request.time` metric, but the `metrics` variable is actually an array, you can easily configure the metric you want to use from an external file or use several metrics to create your custom “node score” function.

Next, you define the scheduler name:

```
scheduler_name = "sysdigsched"
```

This is the name that will be registered on the Kubernetes API, it has to match the pod spec name.

And now, the main loop of the scheduler, it waits for a new event containing an object in `Pending` state and a spec that requires our `scheduler_name`.

```
w = watch.Watch()
for event in w.stream(v1.list_namespaced_pod, "default"):
    if event['object'].status.phase == "Pending" and event['object'].spec.scheduler_name == scheduler_name:
        try:
            print "Scheduling " + event['object'].metadata.name
            res = scheduler(event['object'].metadata.name, best_request_time(nodes_available()))
```

Then it calls the scheduler function, assigning this object (the pending pod) to the node with the current best request time.

How do we measure best request time? With a simple call to the Sysdig API:

```
hostfilter = "host.hostName = '%s'" % hostname
start = -60
end = 0
sampling = 60
metricdata = sdclient.get_data(metrics, start, end, sampling, filter = hostfilter)
```

You will query the metrics we declared before, for the last minute, 60 samples. Then it's just a matter of parsing the data and returning the best value.

A test run with some debugging output will produce an output similar to:

```
Scheduling nginxrc-flk2z
Nodes available: ['kubeworker1', 'kubeworker2']
kubeworker1 (net.http.request.time): 61664.877
kubeworker2 (net.http.request.time): 60456.919
Best node: kubeworker2
```

Here you have the complete scheduler file:

```
1 #!/usr/bin/env python
```

```

4  import random
5  import json
6
7  from kubernetes import client, config, watch
8  from sdcclient import SdcClient
9
10 config.load_kube_config()
11 v1 = client.CoreV1Api()
12 sdclient = SdcClient(<Your Sysdig API token>)
13 sysdig_metric = "net.http.request.time"
14 metrics = [{ "id": sysdig_metric, "aggregations": { "time": "timeAvg", "group": "avg" } }]
15
16 scheduler_name = "sysdigsched"
17
18
19 def get_request_time(hostname):
20     hostfilter = "host.hostName = '%s'" % hostname
21     start = -60
22     end = 0
23     sampling = 60
24     metricdata = sdclient.get_data(metrics, start, end, sampling, filter=hostfilter)
25     request_time = float(metricdata[1].get('data')[0].get('d')[0])
26     print hostname + " (" + sysdig_metric + "): " + str(request_time)
27     return request_time
28
29
30 def best_request_time(nodes):
31     if not nodes:
32         return []
33     node_times = [get_request_time(hostname) for hostname in nodes]
34     best_node = nodes[node_times.index(min(node_times))]
35     print "Best node: " + best_node
36     return best_node
37
38
39 def nodes_available():
40     ready_nodes = []
41     for n in v1.list_node().items:
42         for status in n.status.conditions:
43             if status.status == "True" and status.type == "Ready":
44                 ready_nodes.append(n.metadata.name)
45     return ready_nodes
46
47
48 def scheduler(name, node, namespace="default"):
49     body=client.V1Binding()
50     target=client.V1ObjectReference()
51     target.kind="Node"
52     target.apiVersion="v1"
53     target.name= node
54     meta=client.V1ObjectMeta()
55     meta.name=name
56     body.target=target
57     body.metadata=meta
58     return v1.create_namespaced_binding(namespace, body)
59
60
61 def main():
62     w = watch.Watch()
63     for event in w.stream(v1.list_namespaced_pod, "default"):
64         if event['object'].status.phase == "Pending" and event['object'].spec.scheduler_name == scheduler_name:
65             try:
66                 print "Scheduling " + event['object'].metadata.name
67                 res = scheduler(event['object'].metadata.name, best_request_time(nodes_available()))
68             except client.rest.ApiException as e:
69                 print json.loads(e.body)['message']
70
71
72 if __name__ == '__main__':
73     main()

```

SysdigMonitorKubernetesScheduler.py hosted with ❤ by GitHub

[view raw](#)

Running Kubernetes scheduler in developer mode and live testing

Let's try the script manually first.

Updated Recommendations for You

^

Copy the file to any of your Kubernetes nodes, before you run the script remember to:

- Install the required Python libraries
- Pass your API token to the `sdclient` object
- You may need to adjust the `available_nodes` method using your custom node labels & taints (i.e. remove `NoScheduled` tainted nodes `node-role.kubernetes.io/master:NoSchedule`)

Once you have checked all the items in the list just run it:

```
# python scheduler.py
Nodes available: ['kubeworker1', 'kubeworker2']
kubeworker1 (net.http.request.time): 1202.997
kubeworker2 (net.http.request.time): 1267.912
best node: kubeworker1
Nodes available: ['kubeworker1', 'kubeworker2']
kubeworker1 (net.http.request.time): 1202.997
kubeworker2 (net.http.request.time): 1267.912
best node: kubeworker1
Nodes available: ['kubeworker1', 'kubeworker2']
kubeworker1 (net.http.request.time): 1202.997
kubeworker2 (net.http.request.time): 1267.912
best node: kubeworker1
```

OK, seems that `kubeworker1` had better HTTP response times, so the three `Pending` pods are now running there:

```
$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE
nginx-84cnn   1/1     Running   0           1h    10.244.1.19   kubeworker1
nginx-ffldk   1/1     Running   0           1h    10.244.1.18   kubeworker1
nginx-jq5jk   1/1     Running   0           1h    10.244.1.20   kubeworker1
```

You can use any HTTP load generator to dramatically increase the load in one of the nodes, for example [httpperf](#):

```
kubeworker1:~$ httpperf --server 127.0.0.1 --port 32768 --uri / --num-conn 20000 --num-cal 100000 --rate 200 --timeout 5
```

Scaling the replicationController will automatically generate more pods to be allocated:

```
$ kubectl scale rc nginx --replicas=5
replicationcontroller "nginx" scaled
```

But if we look at the output from our script, `kubeworker2` has now much better responsiveness than `kubeworker1`, currently under `httpperf` stress:

```
Nodes available: ['kubeworker1', 'kubeworker2']
kubeworker1 (net.http.request.time): 17241.543
kubeworker2 (net.http.request.time): 1176.621
best node: kubeworker2
Nodes available: ['kubeworker1', 'kubeworker2']
kubeworker1 (net.http.request.time): 17241.543
kubeworker2 (net.http.request.time): 1176.621
best node: kubeworker2
Your scheduler is allocating new pods in the most responsive node, just as you wanted.
```

Deploy your Kubernetes scheduler as a pod

A better solution is to containerize and orchestrate your scheduler rather than executing a script manually the Kubernetes nodes.

You will need a couple of changes to the script, on line 10:

```
config.load_incluster_config() # instead of config.load_kube_config()
```

And then, on line 12, load the token from a file:

```
sdclient = SdcClient(open("/etc/sysdigtoken/token.txt","r").read().rstrip())
```

From [example-kubernetes-scheduler Github repository](#) you can download some template files that you can use as an starting point:

- `Dockerfile`: to build the scheduler container image
- `scheduler.py`: the modified script file to run as a Docker container in a pod
- `sysdig-account.yaml`: credentials for Kubernetes [RBAC](#)
- `scheduler.yaml`: replicationController to launch the scheduler pod (fill with your container image name)

```
$ kubectl create secret generic sysdig-token --from-file=./token.txt
$ kubectl create -f sysdig-account.yaml
$ kubectl create -f scheduler.yaml
$ kubectl get pods
```

```
pythonscheduler-js8gg    1/1          Running    2           2h
```

Your custom Kubernetes scheduler is ready to go!

Custom Kubernetes scheduler – Golang implementation

During [KubeCon EU 2018](#), we presented a newer and more complete Golang version of the Python code above. You will find the source code and usage instructions [here](#).

This implementation still cannot be considered production ready, however, it has some relevant improvements over the Python version: * Metrics cache & metrics reuse * Failover and failover recovery * Async event handling and scheduling

Further thoughts

This is a relatively simple PoC example, if you really plan to code your own production-level scheduler:

- Declare and properly manage all the possible exception conditions.
- An scheduler has to be fast, benchmark the time it takes to pick a node, average and outliers, maybe you want to use [Sysdig Tracers](#) for that?
- If your code returns an error or is taking too long, you can always code a fallback to the default Kubernetes scheduler, much better than having orphaned pending pods.

A few more use case examples for writing a custom Kubernetes scheduler, we are sure you can come up with your own:

- Schedule backup / storage related pods on nodes with low IO latency and plenty of free HD space.
- Avoid hosts with a high error rate (net error, http error, IO error, etc), you may also set an alarm in Sysdig Monitor.
- Avoid scheduling new pods in a host running container that have specific security incidents. Yes, Sysdig can also look at what your containers are doing from a security point of view with [Sysdig Secure: container run-time security and forensics](#) product.

We hope you found this example useful when diving deep in customizing your Kubernetes cluster behavior. For more deep dives and clear visibility on what your Kubernetes and your containers are doing, check out our [Sysdig Monitor](#) and [Sysdig Secure](#) products and start a free trial yourself.

Post navigation

[Previous Post](#)

[Next Post](#)

Stay up to date

Sign up to receive our newest.

*

Subscription: General Marketing:

☐ Also keep me informed of Sysdig news + updates

Related Posts

Monitoring Kubernetes (part 1)

[Read more](#)

Container Metadata – Understanding Metrics, Labels, & Tags

[Read more](#)

Prometheus metrics / OpenMetrics code instrumentation.

Updated Recommendations for You

^

[Read more](#)
[Start 30-day Free Trial](#)
[Contact Sales](#)

Follow us.



- [Country](#)
 - [日本](#)
 - [United States](#)

Products

- [Sysdig Platform](#)
- [Sysdig Secure](#)
- [Sysdig Monitor](#)

Support

- [Support Log-in](#)
- [Submit Ticket](#)
- [Sysdig Status](#)

Company

- [About Us](#)
- [Leadership](#)
- [Contact Us](#)
- [Sitemap](#)

Copyright 2020 Sysdig, Inc. All Rights Reserved.

[Privacy Policy](#)